

P1.2 Scope + timelines

- **General CFO Agent updates;**

- Data Layer redesign (8 weeks: M1 + M2)
 - Storage + Retrieval updates
 - Rebuild Data layer with facts and dims for data sources (GF/BS/PhL)
 - OLAP workloads
 - Semantic Layer redesign
 - Dedicated KPI library with SQL templates that map to DK and IS
 - Minimize LLM reliance on SQL query generation
 - Data source standardization
 - XLS template for all data sources (CF/BS/PhL)
 - Ingestion pipeline for ETL into CF Agent
 - Built on top of Phase 2 work
- AI Layer redesign (4 weeks: M2)
 - Azure OpenAI API update
 - Chat completions to Responses API
 - Plan + tool calls with one pass
 - Reasoning effort per step
 - Generic Agent workflow updates
 - Redo 5 agents (Orchestrator, SQL, RIA, Graphs, Calculator, Response)
 - Token + Latency optimizations

- **P1.2 Specific changes**

- Domain workbook: Q&A set (4 weeks: M3)
 - Ground truth for 3 tiers of Q&A
 - KPIs required by customer
- MS Fabric migration (4 weeks: M3)
 - SQL database for business data storage
 - Cosmos database for user and application data storage
 - OneLake + Lakehouse for file upload
- UAT Prep (4 weeks: M3+M4)
- Data Ingestion to Fabric (2 weeks: M4)
- Prompt tuning for P1.2 Q&A (3 weeks: M5)
- SIT (3 weeks: M5)
- UAT (2 weeks: M6)

Evaluation only.
Created with Aspose.Slides for Python via .NET 26.2.
Copyright 2004-2026 Aspose Pty Ltd.

Legacy System X-Ray & Navigator

AI-assisted understanding for mainframes and monoliths

From trapped SME knowledge → shared, governed system model → safer modernization



Evaluation only.

Traceable

Created with Aspose.Slides for Python via .NET 26.2.

Client-Controlled

Org-Wide Model

Copyright 2004-2026 Aspose Pty Ltd.



System X-Ray

Parse legacy code into a navigable system map with traceable references



Code Navigator

Production-grade knowledge layer for daily engineering workflows



Modernization

De-risk and accelerate migration using the knowledge layer

The Business Problem



Trapped Knowledge

Critical COBOL / FORTRAN / PL/SQL / Java monoliths run core business, but knowledge is trapped in a few SMEs



Slow Onboarding

Onboarding new engineers and doing impact analysis takes weeks, slows releases, increases risk



Individual, Not Org-Wide

"Just give 5 devs a copilot" helps individuals, but does not create org-wide understanding of the estate

Understanding is the constraint.

Why "Just Give the Team Copilot" Fails

Individual Augmentation

IDE + Copilot Suggestions

- Summaries live in chat history and local context only
- Ephemeral — knowledge is lost when sessions end
- No persistent model: call graphs, data flows, entity maps
- Impact analysis remains SME-dependent

≈ Sticky Notes

Estate Understanding

Persistent System Map

- Persistent system model shared across the entire organization
- Call graphs, data flows, entity maps always up-to-date
- Governed knowledge base for onboarding and impact analysis
- Shared governance eliminates single-point SME dependency


≈ Engineering Atlas

Evaluation only

Created with Aspose.Slides for Python via .NET 26.2.
Copyright 2004-2026 Aspose Pty Ltd.

VS

Our 3-Tier Offering



Tier 1 — Legacy System X-Ray

2–4 weeks



System Map & Impact Analysis

- Parse legacy code, jobs, schemas into navigable system map
- Modules, interfaces, key business entities identified
- Entity-to-code impact maps (Customer / Account / Policy)
- Hot-spot analysis reveals high-risk areas



"Ask Your Legacy System" Demo

- Query flows and business rules with natural language
- Every answer includes traceable code references
- Sources: file path + line numbers for full auditability
- Lightweight demo — value visible in days, not months

DELIVERABLES

System Atlas

Impact Maps

Q&A Demo

Findings Readout

Tier 2 — Legacy Code Navigator

Team Copilot



Impact Analysis Trace all code affected by a proposed change. Sources: file path + line #	Business Flow Tracing Follow a business process through the codebase. Sources: file path + line #	Explain This Program Get a clear summary of what any module does. Sources: file path + line #	Data Lineage Track data from source to consumption point. Sources: file path + line #
--	--	--	--

GUIDED ONBOARDING

1. System Overview


2. Key Flows

3. Deep Dive

Deployed in client VPC • Optional IDE / web integration

Tier 3 — Modernization Accelerator

STRANGLER PATTERN



Gradually hollow out the monolith → route traffic → migrate to services

GENERATED ARTIFACTS

Interface Specs

Candidate API contracts generated from existing flows

Sequence Diagrams

Visualize interactions between components

Test Suggestions

Ensure new code preserves legacy semantics

AI-assisted refactoring preserves legacy semantics while moving to modern stacks



How This Differs from Tool Licenses

Dimension	Tool License	X-Ray & Navigator
Knowledge Model	Ad-hoc file summaries that are ephemeral and individual	Persistent system model: call graphs, data flows, entity maps
Governance	Individual use with no shared context across users	Shared across teams: onboarding tours, impact reports, dashboards
Deployment	Cloud-hosted, vendor-controlled infrastructure	Client environment (on-prem / VPC) with full audit trail

Evaluation only.
Created with Aspose.Slides for Python via .NET 26.2.
Copyright 2004-2026 Aspose Pty Ltd.

NON-NEGOTIABLES



Traceable

Auditable

VPC / On-Prem

Business Outcomes

Months → **Weeks**

Onboarding Time
on Legacy Systems

Tailor to client baseline

Days → **Hours**

Impact Analysis
Effort for Changes

Tailor to client baseline

High → **Lower**

Dependency on
Legacy SMEs

Tailor to client baseline


KNOWLEDGE FLYWHEEL

Learn → Document → Query → Improve → Modernize

Create a living knowledge base for future modernization

Engagement Approach

Typical Timeline



RISK CONTROLS

Data Residency

Access Control

Audit Logs

DELIVERABLES — STARTER KIT

System Atlas

Impact Maps

Query Recipes

Deployment Plan